

AI Drop-Service OS Agent Training

Generated 2026-07-03T07:28:46Z from team board and persistent-agent packet

Bottom line: Research Scout and QA Red-Team now exist as persistent isolated OpenClaw agents. Use them as bounded departments. Veritas main remains the orchestrator, verifier, and final truth surface.

OWNER WORKFLOW

WF75

AI Drop-Service OS

AGENTS CREATED


2

zero external bindings

STATUS

phase1_internal_agents_created_training_complete

ARCHITECTURE



AGENT PROFILES

Agent	Model	Role	Best first use
research-scout Research Scout	openai/gpt-5.4-mini	public-source research and source tables	Ask it for public examples, competitor scans, tool/vendor maps, and source-grounded market evidence.
qa-redteam QA Red-Team	openai/gpt-5.4	independent critique, risk review, acceptance criteria, proof gaps	Ask it to challenge claims, offers, SOPs, prompts, and workflow artifacts before trusting them.

How You Communicate With Agents

Option 1

Ask Veritas to route
safest daily path

Option 2

Use CLI agent turn
direct, internal, no channel binding

Option 3

Persistent session key
good for continuing a lane

Later

External channels or cron
owner-gated, not enabled

SAFEST DAILY PATH

Talk to Veritas main and ask me to route the task. This gives you agent leverage without losing integration, memory, and QA control.

Veritas, send this to research-scout:

Find 8 sourced examples of AI workflow audit or AI automation sprint offers for small service businesses. Return source links, pricing if available, promise, deliverables, and risk notes.

Veritas, send this to qa-redteam:

Challenge the AI Workflow Clarity Sprint offer. Find the top 10 claims, privacy, delivery, pricing, and proof risks. Return blockers first.

DIRECT CLI PATH

Use this when you want to message a configured isolated agent directly through OpenClaw Gateway.

```
openclaw agent --agent research-scout --message "Find 8 sourced examples of AI workflow audit offers. Return a source table and cite URLs." --json
```

```
openclaw agent --agent qa-redteam --message "Review this offer for weak claims, privacy risk, delivery risk, and missing proof. Return severity-ranked findings." --json
```

Current WebChat limitation: direct `@research-scout` or `@qa-redteam` style routing is not enabled in this chat. OpenClaw blocked the agent-to-agent session-send tool until cross-agent visibility is enabled in config. I did not change that config in this slice. Use the CLI path, or ask Veritas main to route the task for you.

PERSISTENT SESSION KEY

Use a stable session key when a line of work should continue across turns.

```
openclaw agent --agent research-scout --session-key agent:research-scout:wf75-research --message "Continue the competitor scan. Add 5 examples and dedupe against prior notes." --json
```

```
openclaw sessions --agent research-scout --json
openclaw sessions --agent qa-redteam --json
```

Your Role, Responsibilities, And Prompt Recipes

RANDALL OWNS

- Objective and priority.
- Truth about constraints, budget, customer access, and risk tolerance.
- Final approval for external channels, cron schedules, customer use, vendors, payments, domains, and outreach.
- Decision on whether an output becomes an offer, demo, or build lane.

VERITAS OWNS

- Routing work to the right agent.
- Checking helper output against source artifacts.
- Integrating final answers and preserving memory.
- Stopping work when authority, privacy, proof, or external-action boundaries are unclear.

PROMPT RECIPES

Use case	Prompt shape
Research Scout	Find [number] sourced examples of [market/offer/tool]. Return source URL, buyer, promise, price if public, deliverables, evidence quality, and what we can learn for WF75. Do not contact anyone.
QA Red-Team	Review [artifact/offer/workflow]. Return blockers first, severity-ranked risks, missing proof, privacy concerns, weak claims, and acceptance criteria. No edits unless assigned.
Multi-agent chain	Veritas, send Research Scout this research task, then send the result to QA Red-Team for critique, then integrate both into a recommendation.

TREND RADAR

- Vertical AI services are moving from generic chatbot promises to specific workflow outcomes.
- SMB missed-lead intake, follow-up rescue, SOP cleanup, support triage, and weekly research briefs are practical first offers.
- Buyers care about trust, privacy, proof, and ROI more than the word AI.
- Multi-agent systems need task ownership, memory, stop lines, and QA gates.
- Human review is part of the moat because it prevents polished but unreliable output.

Stop Lines And Next Safe Exercise

HARD STOP LINES

- Do not bind agents to external channels without separate owner approval.
- Do not let helper output become final truth without Veritas verification.
- Do not give agents credentials, customer systems, vendor accounts, payment accounts, ads, domains, or external delivery authority in this slice.
- Do not infer finance, portfolio, paper/live trade, account, or capital authority.
- Do not schedule autonomous cron dispatch without a separate cron contract and owner approval.

Current authority: internal training and bounded internal agent use only. No Telegram, Discord, email, webhook, customer-facing routing, cron dispatch, real customer data, vendor/payment setup, ads, domains, credentials, finance/account action, or external delivery has been approved or enabled.

FIRST TRAINING EXERCISE

1. Ask Research Scout for a source table on missed-lead automation offers for local service businesses.
2. Ask Veritas to summarize what is usable and what is weak.
3. Send the summary to QA Red-Team for blockers and acceptance criteria.
4. Have Veritas integrate both into a fictional demo brief for the AI Workflow Clarity Sprint.

WHAT SUCCESS LOOKS LIKE

- Research contains real sources and a clear source-quality column.
- QA catches hype, privacy, delivery, pricing, and proof problems.
- Veritas produces the final answer and records the durable next action.
- No external action happens until you explicitly approve that separate step.